Python Project Mall

May 5, 2023

[1]:
```
import pandas as pd # Data manipulation
import seaborn as sns # Statistical visualization library
import matplotlib.pyplot as plt # Another visualization library
from sklearn.cluster import KMeans # For create clusters
import warnings
warnings.filterwarnings("ignore")
```

[2]:
```
df = pd.read_csv("C:/Users/Drac_/OneDrive/Desktop/Python_Project_Data/
↪Mall_Customers.csv")
```

[3]:
```
df.head()
```

[3]:

	CustomerID	Gender	Age	Annual Income (k$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

1 Univariate Analysis

[4]:
```
df.describe()
```

[4]:

	CustomerID	Age	Annual Income (k$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

[5]:
```
sns.displot(df["Annual Income (k$)"])
```

[5]:
```
<seaborn.axisgrid.FacetGrid at 0x2146169b0a0>
```

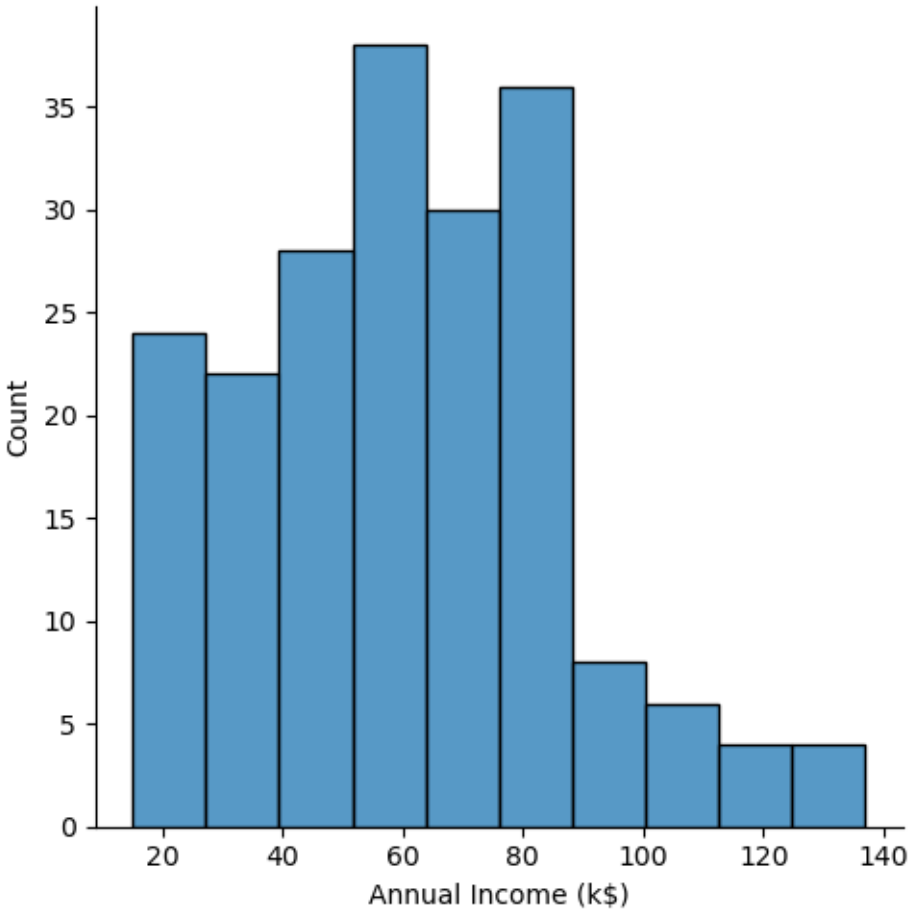

[6]:
```
df.columns
```

[6]:
```
Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k$)',
       'Spending Score (1-100)'],
      dtype='object')
```

[7]:
```
columns = ['Age', 'Annual Income (k$)','Spending Score (1-100)']
for i in columns:
    plt.figure()
    sns.displot(df[i])
```

```
<Figure size 640x480 with 0 Axes>
```

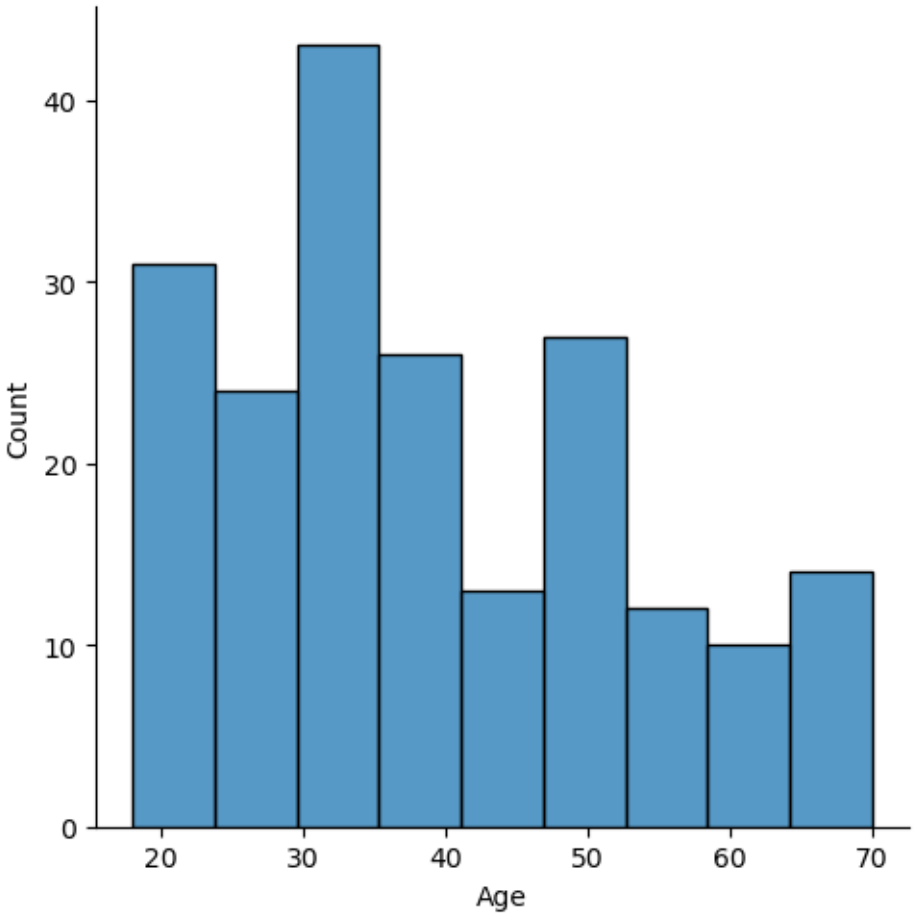


```
<Figure size 640x480 with 0 Axes>
```

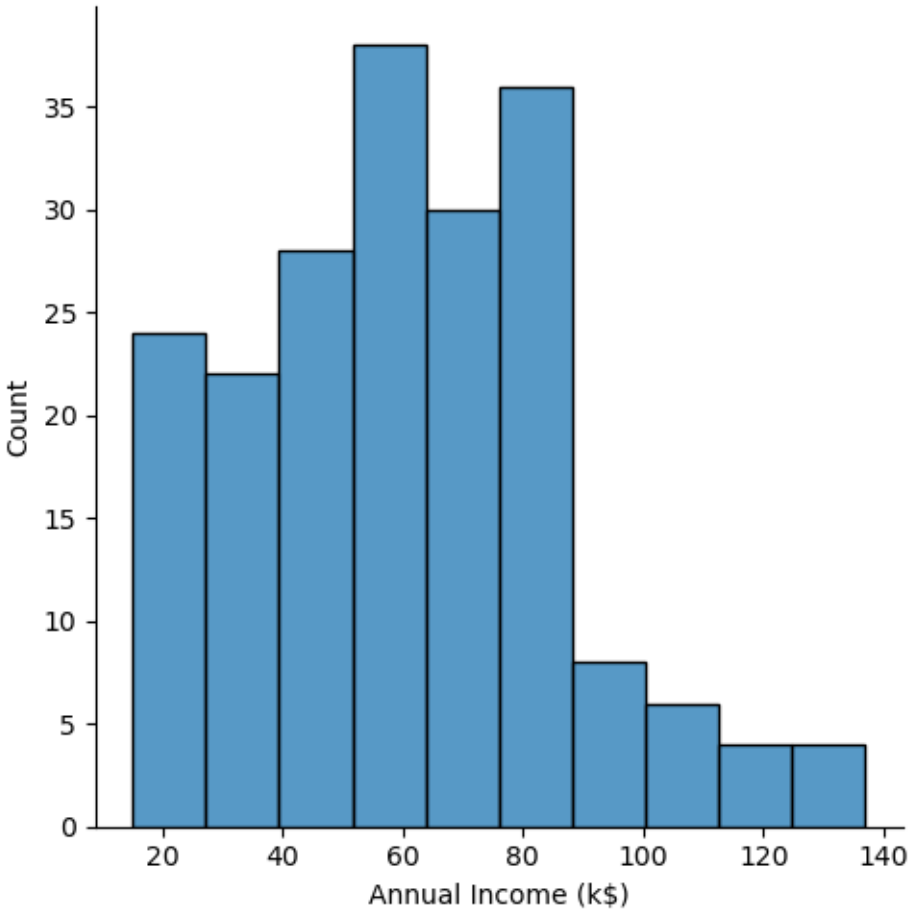


```
<Figure size 640x480 with 0 Axes>
```

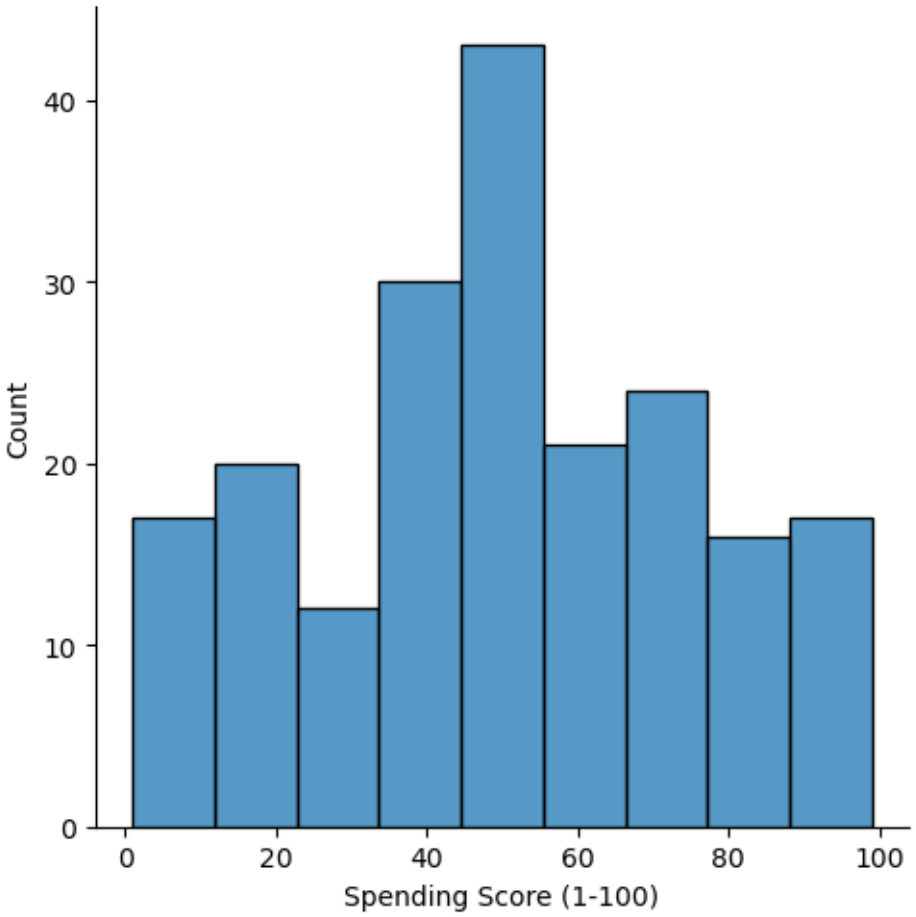

[8]:
```
sns.kdeplot(x=df["Annual Income (k$)"],shade=True,hue=df["Gender"]);
```

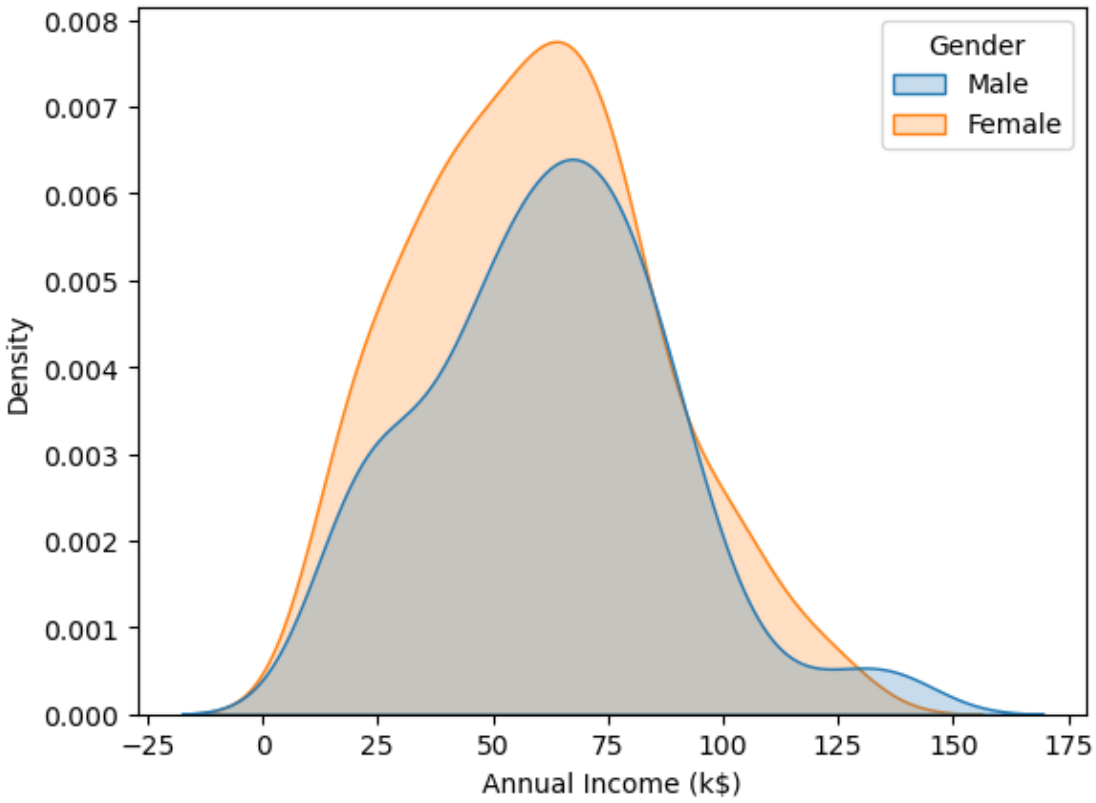

[9]:
```
columns = ['Age', 'Annual Income (k$)','Spending Score (1-100)']
for i in columns:
    plt.figure()
    sns.kdeplot(x=df[i],shade=True,hue=df["Gender"]);
```

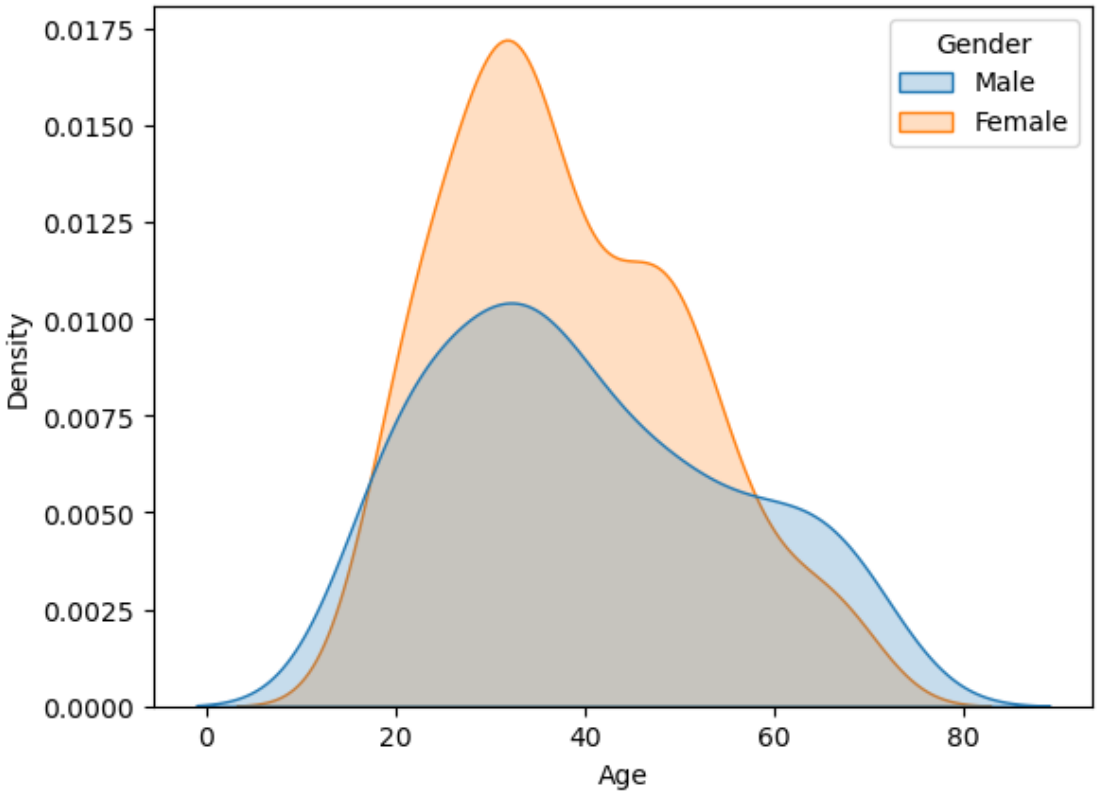
Gender
Male
Female
Density
0.0175
0.0150
0.0125
0.0100
0.0075
0.0050
0.0025
0.0000
0
20
40
60
80
Age

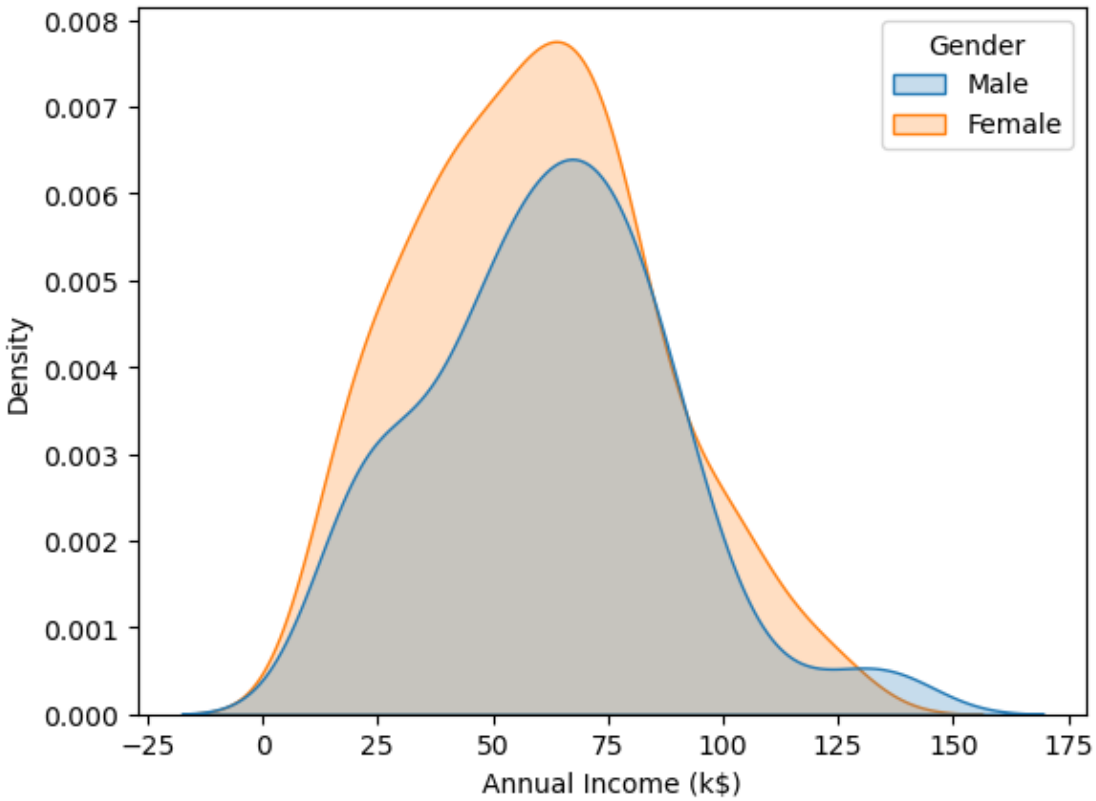
Gender
Male
Female
Density
0.008
0.007
0.006
0.005
0.004
0.003
0.002
0.001
0.000
−25
0
25
50
75
100
125
150
175
Annual Income (k$)

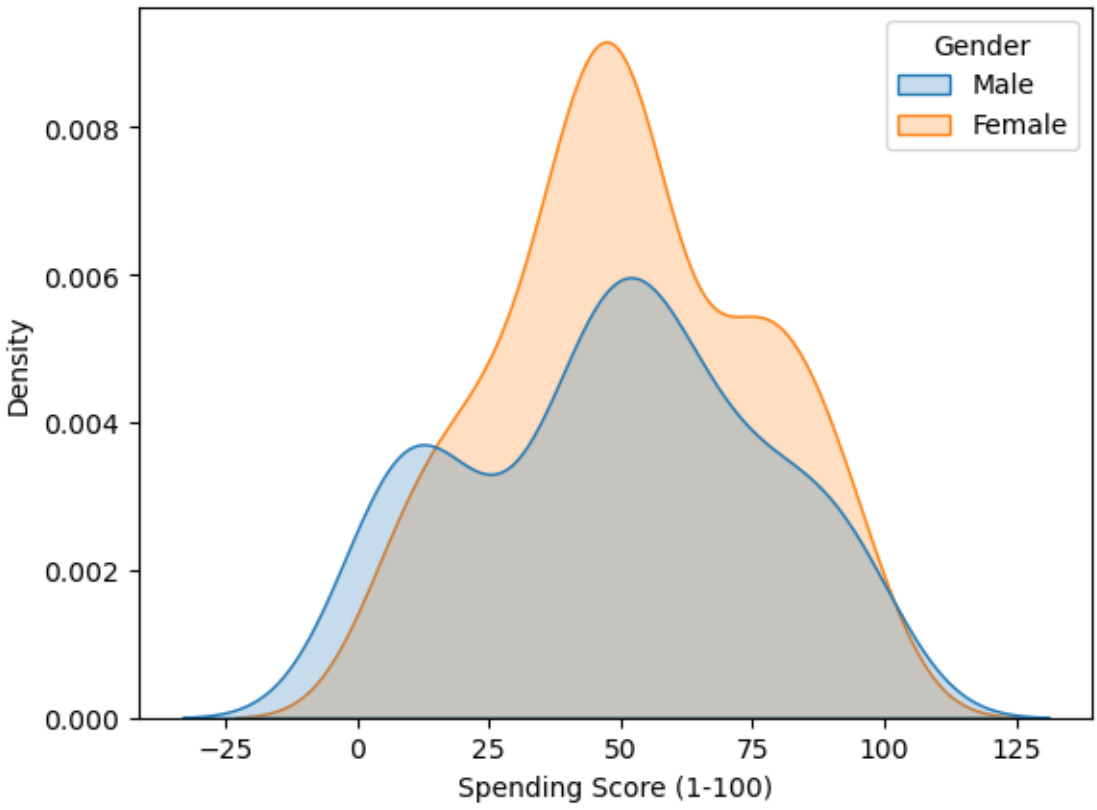

[10]:
```
columns = ['Age', 'Annual Income (k$)','Spending Score (1-100)']
for i in columns:
    plt.figure()
    sns.boxplot(data=df,x="Gender",y=df[i]);
```

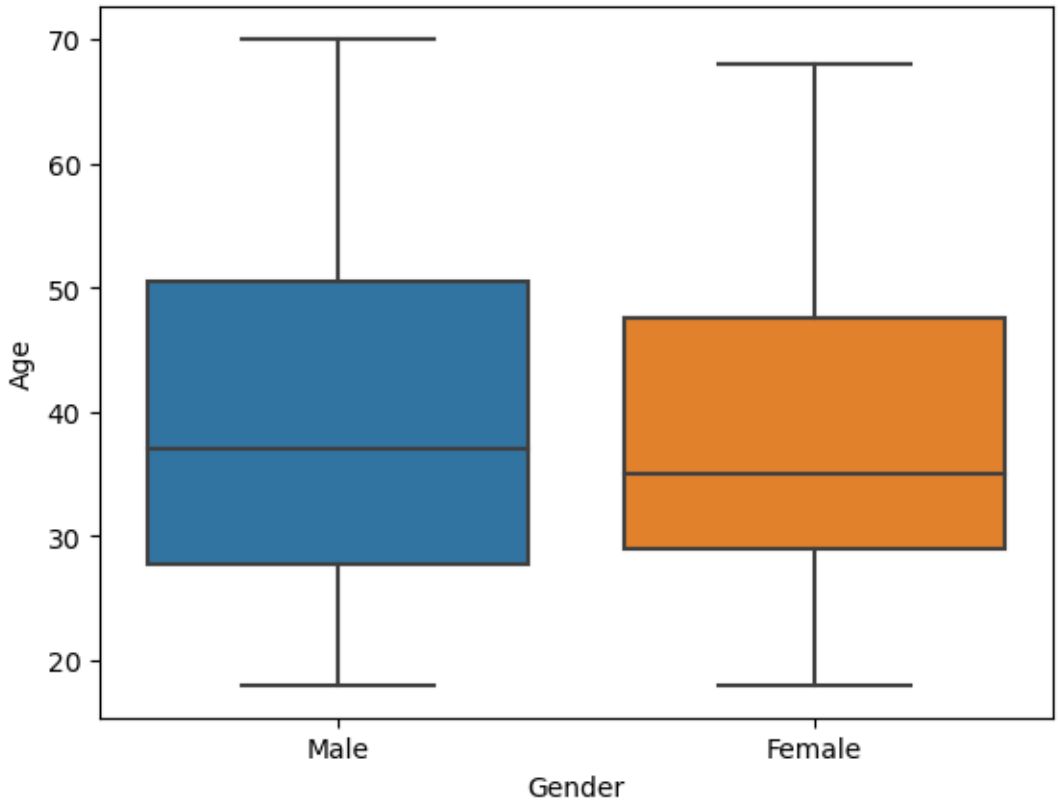
70
60
50
40
30
20
Age
Male
Female
Gender

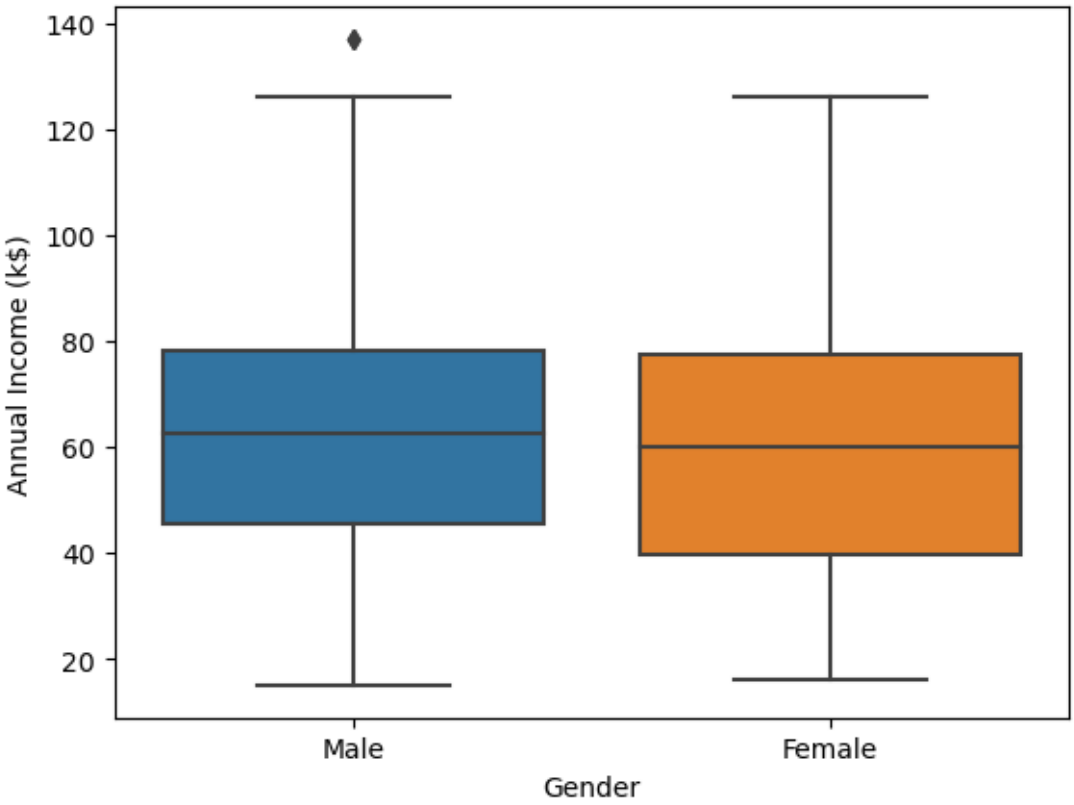
140
120
100
80
60
40
20
Annual Income (k$)
Male
Female
Gender

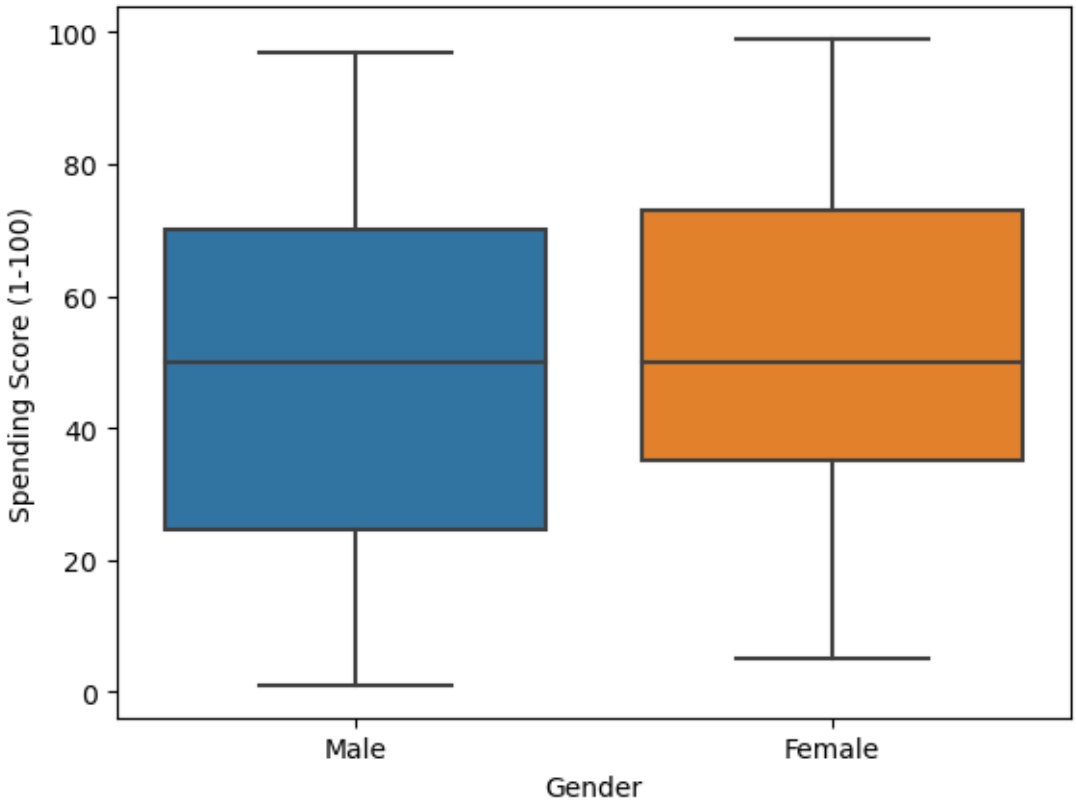

[11]:
```
df["Gender"].value_counts(normalize=True)
```

[11]:
```
Female    0.56
Male      0.44
Name: Gender, dtype: float64
```

2 Bivariate Analysis

[12]:
```
sns.scatterplot(data=df, x="Annual Income (k$)", y='Spending Score (1-100)')
```

[12]:
```
<Axes: xlabel='Annual Income (k$)', ylabel='Spending Score (1-100)'>
```

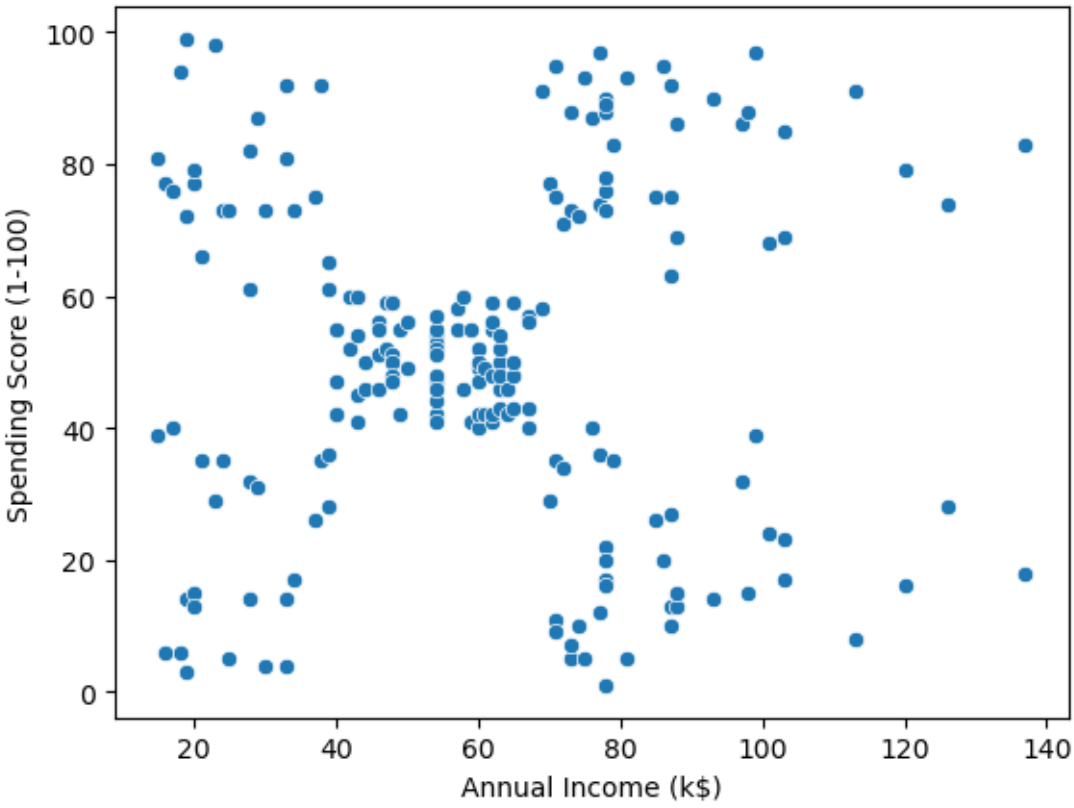

[13]:
```
#df=df.drop("CustomerID",axis=1) <- Don't need to run this again
sns.pairplot(df, hue="Gender")
```

[13]: <seaborn.axisgrid.PairGrid at 0x21462f8e620>

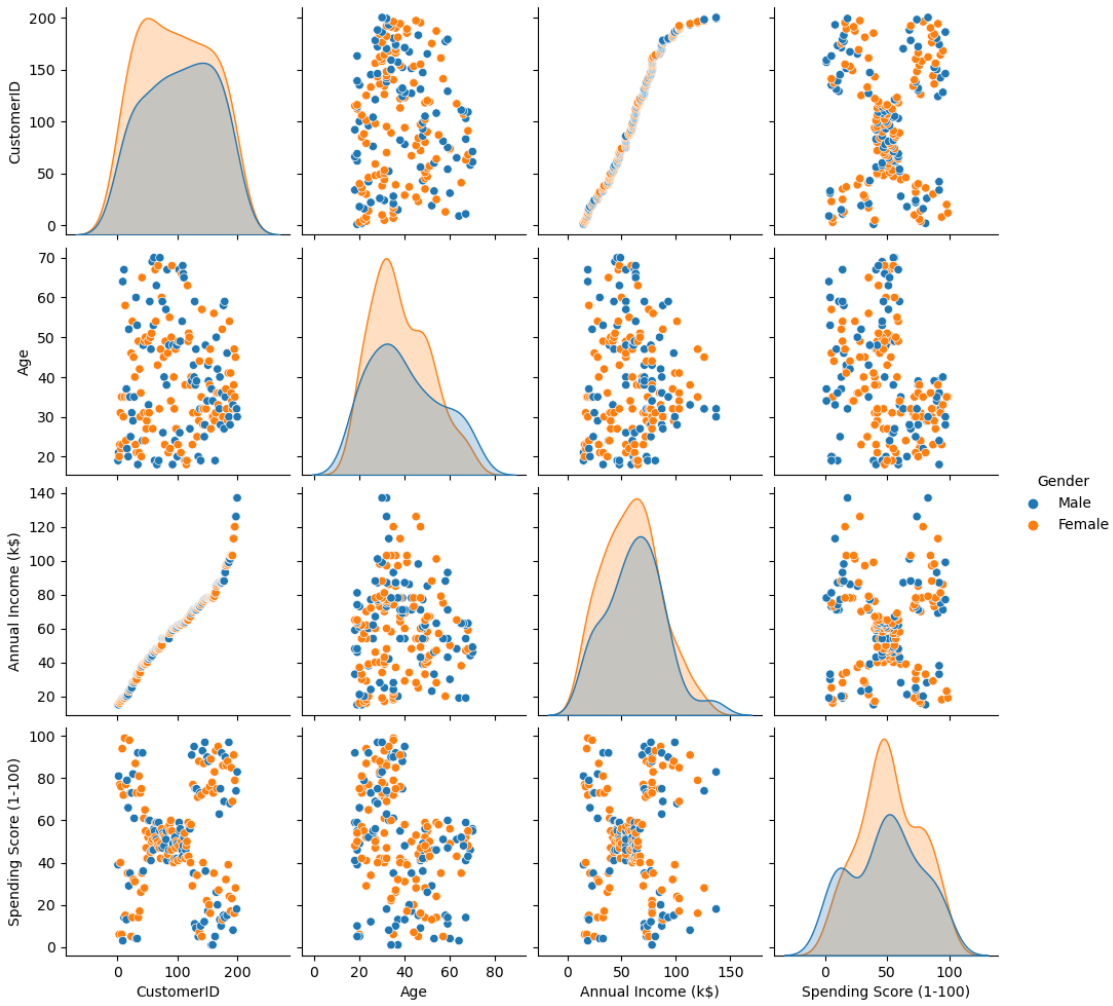

[14]:
```
df.groupby(["Gender"])["Age", "Annual Income (k$)",
       "Spending Score (1-100)"].mean()
```

[14]:

Gender	Age	Annual Income (k$)	Spending Score (1-100)
Female	38.098214	59.250000	51.526786
Male	39.806818	62.227273	48.511364

[15]:
```
df.corr() # <- Correlation function
```

[15]:

	CustomerID	Age	Annual Income (k$)	\
CustomerID	1.000000	-0.026763	0.977548	
Age	-0.026763	1.000000	-0.012398	
Annual Income (k$)	0.977548	-0.012398	1.000000	
Spending Score (1-100)	0.013835	-0.327227	0.009903	

	Spending Score (1-100)
CustomerID	0.013835
Age	-0.327227
Annual Income (k$)	0.009903
Spending Score (1-100)	1.000000

[16]:
```
sns.heatmap(df.corr(), annot=True, cmap="coolwarm")
```

[16]: <Axes: >

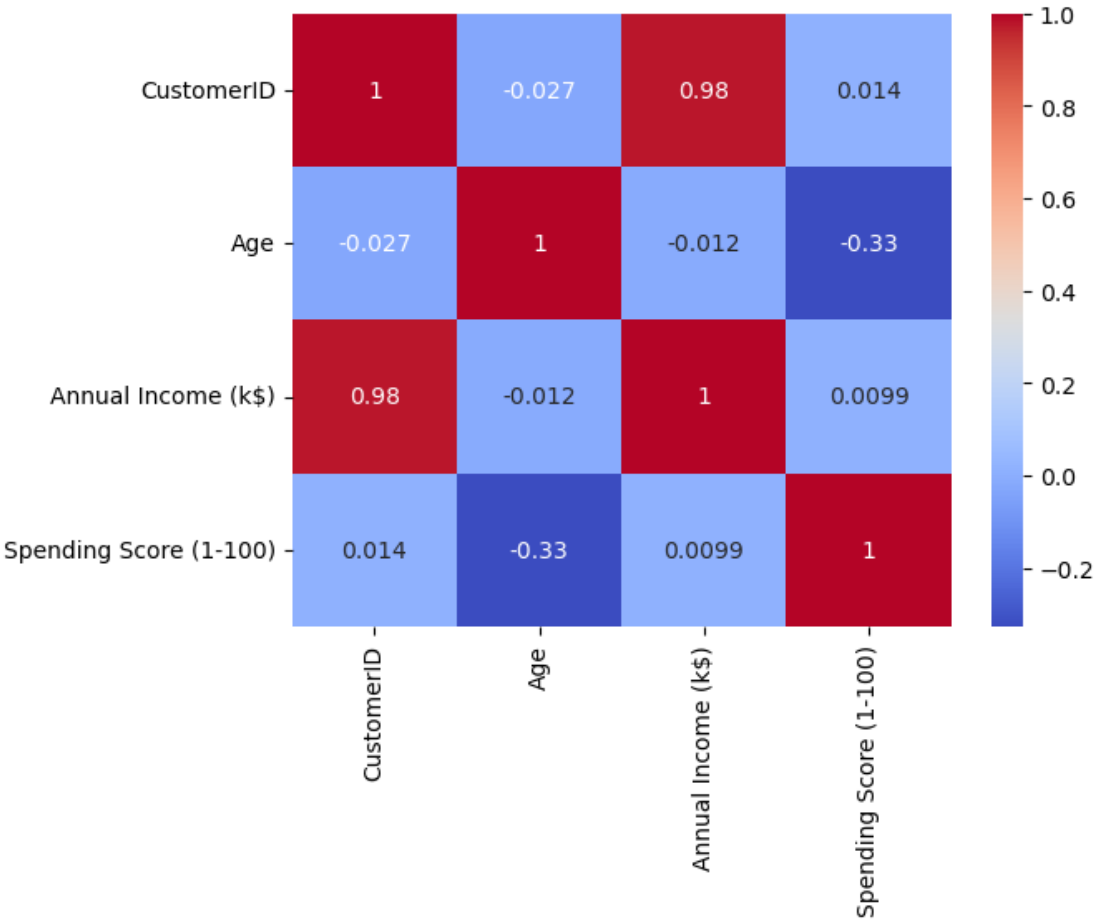

3 Clustering - Univariate, Bivariate, Multivariate

[68]:
```
clustering1 = KMeans(n_clusters=3)
```

[69]:
```
clustering1.fit(df[["Annual Income (k$)"]])
```

[69]: KMeans(n_clusters=3)

[70]:
```
clustering1.labels_
```

[70]:
```
array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2])
```

[71]:
```
df["Income Cluster"] = clustering1.labels_
df.head()
```

[71]:

	CustomerID	Gender	Age	Annual Income (k$)	Spending Score (1-100)	Income Cluster
0	1	Male	19	15	39	1
1	2	Male	21	15	81	1
2	3	Female	20	16	6	1
3	4	Female	23	16	77	1
4	5	Female	31	17	40	1

[72]:
```
df["Income Cluster"].value_counts()
```

[72]:
```
0    90
1    74
2    36
Name: Income Cluster, dtype: int64
```

[73]:
```
clustering1.inertia_ #Inertia represents is the distance between centroids
```

[73]:
```
23517.33093093093
```

[74]:
```
intertia_scores=[]
for i in range(1,11):
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(df[["Annual Income (k$)"]])
    intertia_scores.append(kmeans.inertia_)
```

[75]:
```
intertia_scores
```

[75]:
```
[137277.28000000003,
 48660.88888888888,
 23517.33093093093,
 13278.112713472487,
 8481.496190476191,
 5050.9047619047615,
 3976.358363858364,
 2822.499694749695,
 2173.287445887446,
 1859.0235042735042]
```

[76]:
```
plt.plot(range(1,11),intertia_scores)
```

[76]:
```
[<matplotlib.lines.Line2D at 0x2146930a800>]
```

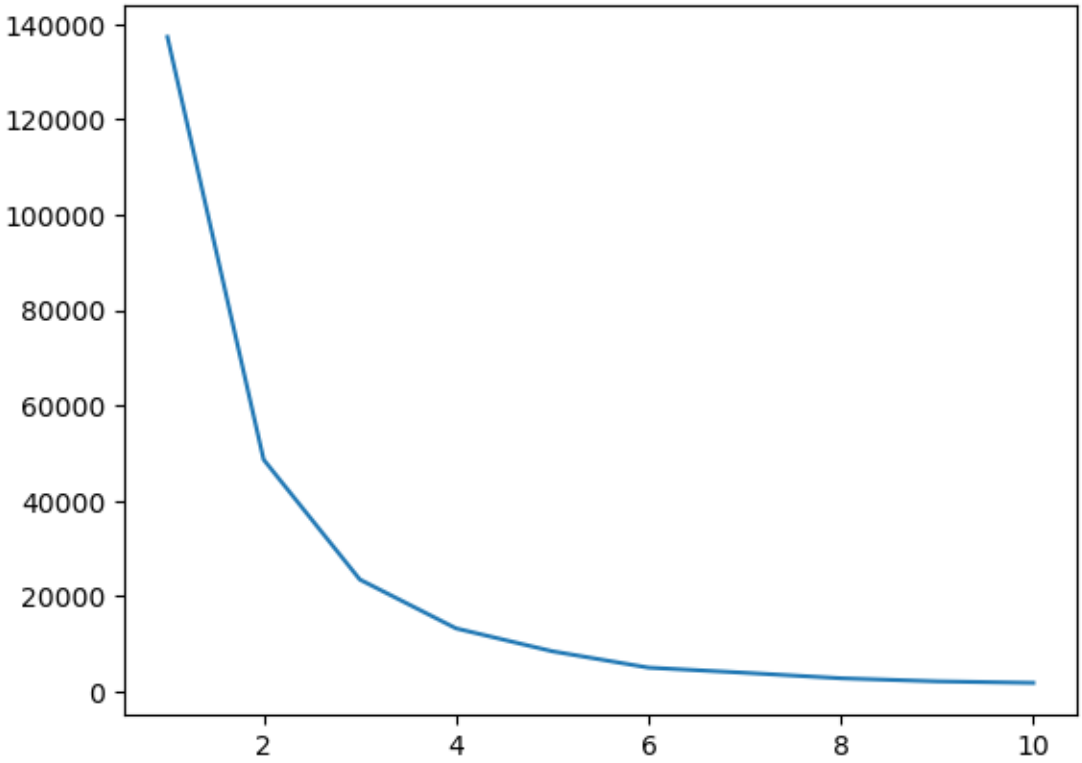

[77]:
```
df.columns
```

[77]:
```
Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k$)',
       'Spending Score (1-100)', 'Income Cluster'],
      dtype='object')
```

[78]:
```
df.groupby("Income Cluster")['Age', 'Annual Income (k$)', 'Spending Score␣
↪(1-100)'].mean()
```

[78]:

Income Cluster	Age	Annual Income (k$)	Spending Score (1-100)
0	38.722222	67.088889	50.000000
1	39.500000	33.486486	50.229730
2	37.833333	99.888889	50.638889

[79]:
```
#Bivariate Clustering
```

[84]:
```
clustering2 = KMeans(n_clusters=5)
clustering2.fit(df[['Annual Income (k$)','Spending Score (1-100)']])
df["Spending and Income Cluster"] = clustering2.labels_
df.head()
```

[84]:

	CustomerID	Gender	Age	Annual Income (k$)	Spending Score (1-100)	\
0	1	Male	19	15	39	
1	2	Male	21	15	81	
2	3	Female	20	16	6	
3	4	Female	23	16	77	
4	5	Female	31	17	40	

	Income Cluster	Spending and Income Cluster
0	1	3
1	1	1
2	1	3
3	1	1
4	1	3

[85]:
```
intertia_scores2=[]
for i in range(1,11):
    kmeans2 = KMeans(n_clusters=i)
    kmeans2.fit(df[["Annual Income (k$)","Spending Score (1-100)"]])
    intertia_scores2.append(kmeans2.inertia_)
plt.plot(range(1,11),intertia_scores2)
```

[85]:
```
[<matplotlib.lines.Line2D at 0x214694793f0>]
```

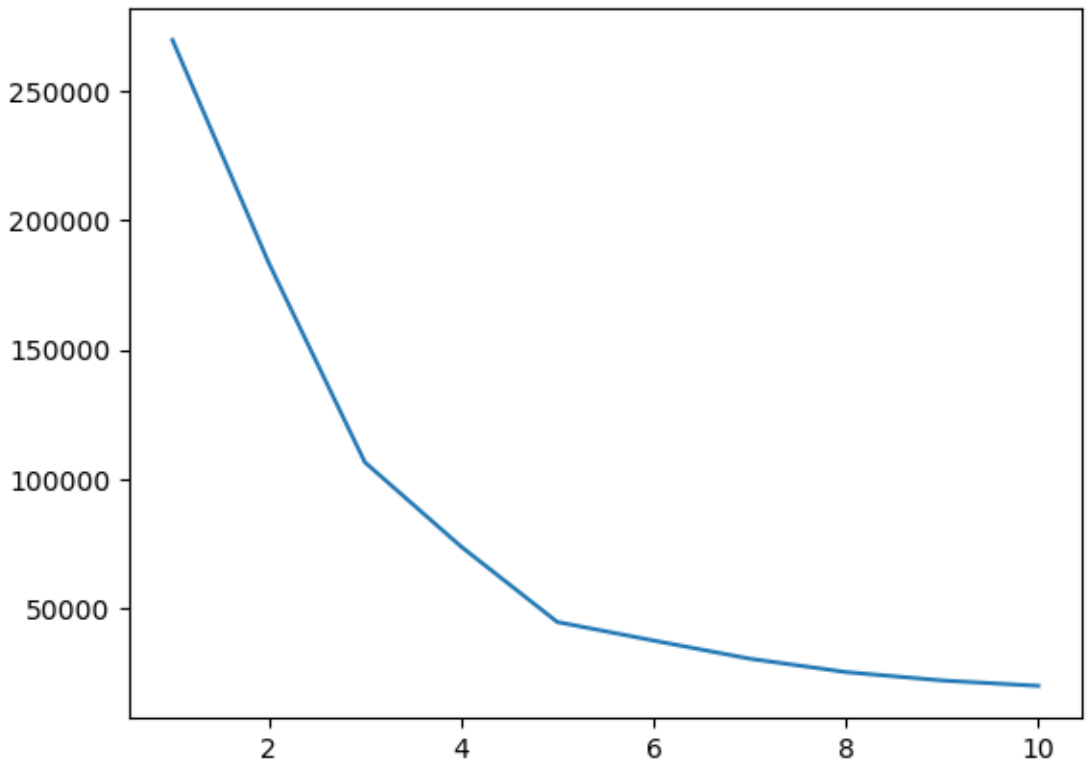

[100]:
```
centers = pd.DataFrame(clustering2.cluster_centers_)
centers.columns = ["x","y"]
```

[116]:
```
plt.figure(figsize = (10,8))
plt.scatter(x=centers["x"], y=centers["y"],c="black",marker="*")
sns.scatterplot(data=df, x = "Annual Income (k$)", y = "Spending Score␣
↪(1-100)", hue = "Spending and Income Cluster", palette = "tab10")
plt.savefig("Clustering_bivariate.png")
```

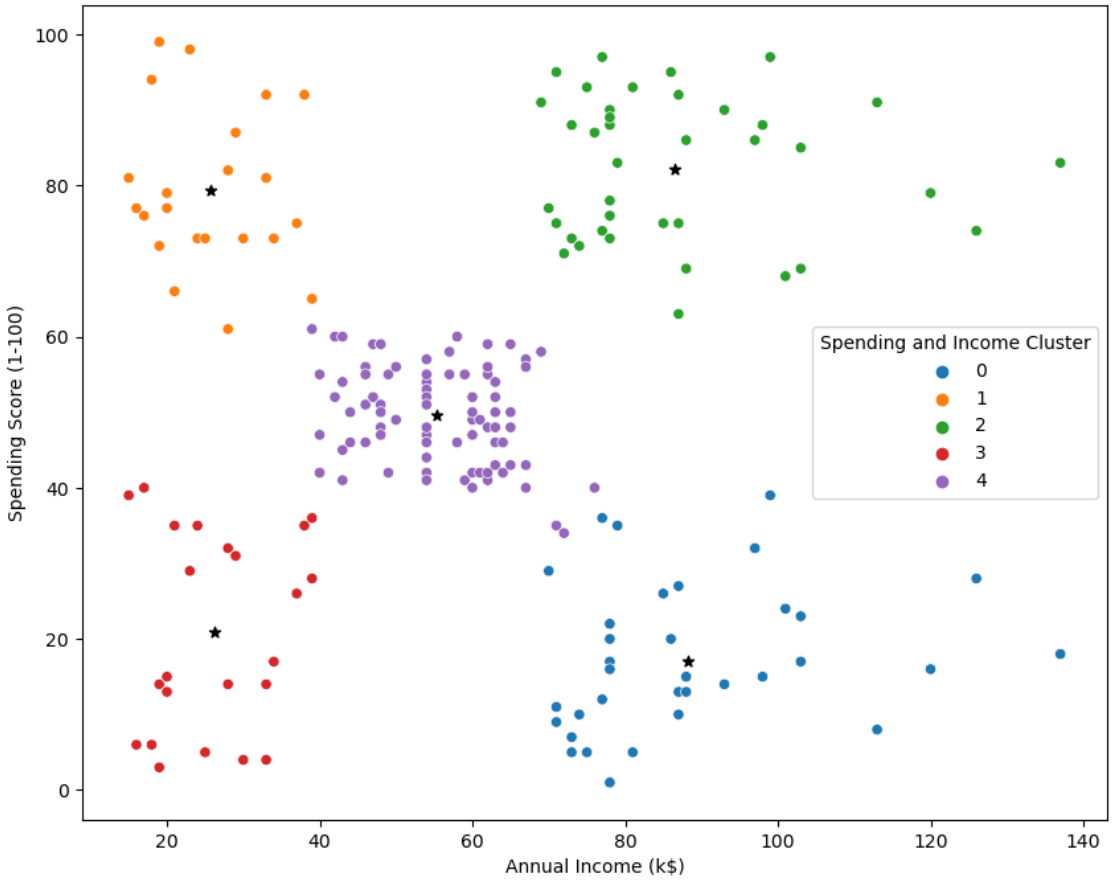

[104]:
```
pd.crosstab(df["Spending and Income Cluster"],df["Gender"], normalize="index")
```

[104]:

Gender	Female	Male
Spending and Income Cluster		
0	0.457143	0.542857
1	0.590909	0.409091
2	0.538462	0.461538
3	0.608696	0.391304
4	0.592593	0.407407

[105]:
```
df.groupby("Spending and Income Cluster")["Age","Annual Income (k$)",
       "Spending Score (1-100)"].mean()
```

[105]:

	Age	Annual Income (k$)	\
Spending and Income Cluster			
0	41.114286	88.200000	
1	25.272727	25.727273	
2	32.692308	86.538462	
3	45.217391	26.304348	

```
4                              42.716049           55.296296

                               Spending Score (1-100)
Spending and Income Cluster
0                                          17.114286
1                                          79.363636
2                                          82.128205
3                                          20.913043
4                                          49.518519
```

[106]:
```
#multivariate clustering
from sklearn.preprocessing import StandardScaler
```

[107]:
```
scale = StandardScaler()
```

[108]:
```
df.head()
```

[108]:
```
   CustomerID  Gender  Age  Annual Income (k$)  Spending Score (1-100)  \
0           1    Male   19                  15                      39
1           2    Male   21                  15                      81
2           3  Female   20                  16                       6
3           4  Female   23                  16                      77
4           5  Female   31                  17                      40

   Income Cluster  Spending and Income Cluster
0               1                            3
1               1                            1
2               1                            3
3               1                            1
4               1                            3
```

[109]:
```
#dff = pd.get_dummies(df) <- return values from female and male, but just we␣
 ↪need one value, use drop for that
dff = pd.get_dummies(df,drop_first=True)
dff.head()
```

[109]:
```
   CustomerID  Age  Annual Income (k$)  Spending Score (1-100)  \
0           1   19                  15                      39
1           2   21                  15                      81
2           3   20                  16                       6
3           4   23                  16                      77
4           5   31                  17                      40

   Income Cluster  Spending and Income Cluster  Gender_Male
0               1                            3            1
1               1                            1            1
2               1                            3            0
```

```
3                1                              1              0
4                1                              3              0
```

[110]:
```
dff.columns
```

[110]:
```
Index(['CustomerID', 'Age', 'Annual Income (k$)', 'Spending Score (1-100)',
       'Income Cluster', 'Spending and Income Cluster', 'Gender_Male'],
      dtype='object')
```

[111]:
```
dff = dff[['Age', 'Annual Income (k$)', 'Spending Score (1-100)','Gender_Male']]
dff.head()
```

[111]:

	Age	Annual Income (k$)	Spending Score (1-100)	Gender_Male
0	19	15	39	1
1	21	15	81	1
2	20	16	6	0
3	23	16	77	0
4	31	17	40	0

[112]:
```
dff = pd.DataFrame(scale.fit_transform(dff))
dff.head()
```

[112]:

	0	1	2	3
0	-1.424569	-1.738999	-0.434801	1.128152
1	-1.281035	-1.738999	1.195704	1.128152
2	-1.352802	-1.700830	-1.715913	-0.886405
3	-1.137502	-1.700830	1.040418	-0.886405
4	-0.563369	-1.662660	-0.395980	-0.886405

[113]:
```
intertia_scores3=[]
for i in range(1,11):
    kmeans3 = KMeans(n_clusters=i)
    kmeans3.fit(dff)
    intertia_scores3.append(kmeans3.inertia_)
plt.plot(range(1,11),intertia_scores3)
```

[113]:
```
[<matplotlib.lines.Line2D at 0x2146bb31a50>]
```

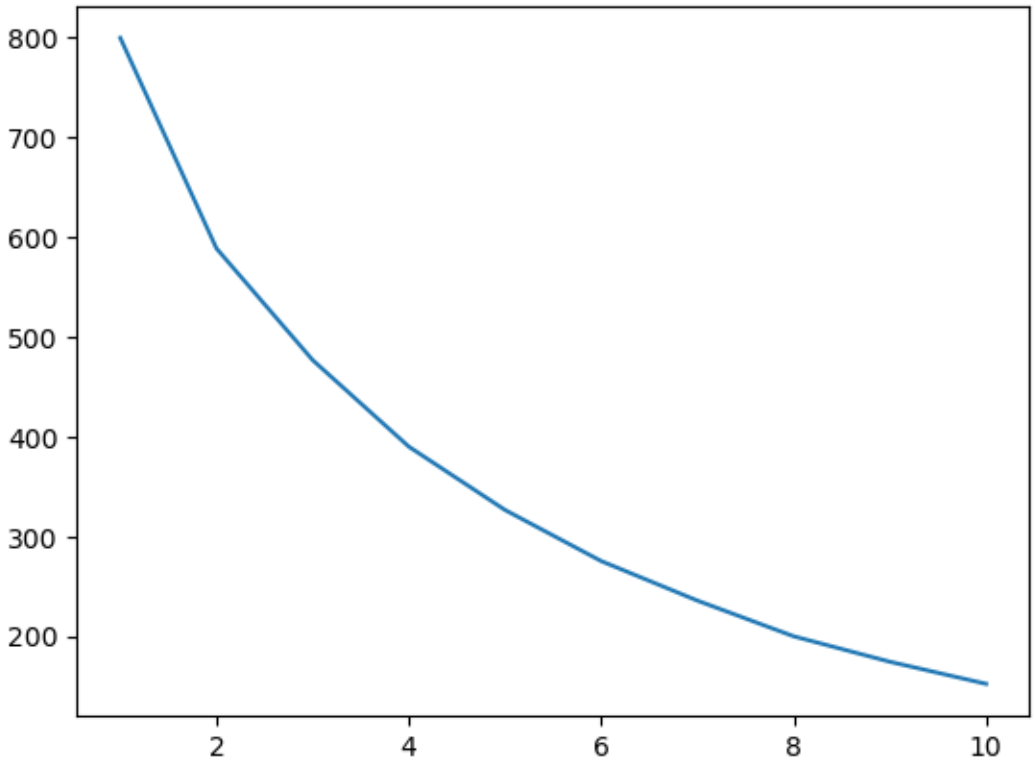

[115]:
```
df
```

[115]:
```
     CustomerID  Gender  Age  Annual Income (k$)  Spending Score (1-100)  \
0             1    Male   19                  15                      39
1             2    Male   21                  15                      81
2             3  Female   20                  16                       6
3             4  Female   23                  16                      77
4             5  Female   31                  17                      40
..          ...     ...  ...                 ...                     ...
195         196  Female   35                 120                      79
196         197  Female   45                 126                      28
197         198    Male   32                 126                      74
198         199    Male   32                 137                      18
199         200    Male   30                 137                      83

     Income Cluster  Spending and Income Cluster
0                 1                            3
1                 1                            1
2                 1                            3
3                 1                            1
4                 1                            3
..              ...                          ...
```

```
195              2                       2
196              2                       0
197              2                       2
198              2                       0
199              2                       2

[200 rows x 7 columns]
```

[117]:
```
df.to_csv("Clustering.csv")
```

[ ]:
```

```